

# Dermatologist Review & Recommendation API

## 1. Architecture Overview

The system follows a layered architecture pattern with clear separation of concerns. To view the system architecture kindly visit this [link](#).

### Core Components:

- 1.1. API Layer:** FastAPI handles requests, routing and response formatting
- 1.2. Authentication:** JWT-based token system for secure access control
- 1.3. Data Layer:** SQLAlchemy ORM with SQLite for data persistence
- 1.4. Validation:** Pydantic models ensure data integrity and API contract compliance
- 1.5. External Integration:** HTTP client for fetching product data from DummyJSON API

## 2. Data Models and Relationships

To view the ER Diagram, please visit this [link](#).

### 2.1 Entity Definitions:

- **User:** Base entity for authentication (patients and doctors)
- **Doctor:** Extended profile for medical practitioners with specialization
- **Rating:** Numerical rating (1-5) linking users to doctors
- **Review:** Text feedback with word limit validation
- **Recommendation:** Product suggestions with UUID-based sharing and expiry

### 2.2 Key Relationships:

- One-to-many: Doctor → Ratings, Reviews, Recommendations
- Many-to-one: Ratings, Reviews → User
- One-to-one: User → Doctor (for medical practitioners)

## 3. Key Design Decisions

### 3.1. Authentication Strategy

**Decision:** JWT token-based authentication with role differentiation

#### Rationale:

- Stateless authentication suitable for REST APIs
- Built-in expiry mechanism
- Role-based access control for doctor-specific operations

### 3.2. Database Choice

**Decision:** SQLite in-memory with SQLAlchemy ORM

#### Rationale:

- Meets assessment requirement for in-memory storage
- SQLAlchemy provides database abstraction for future migration
- Relational model suits the structured data relationships

### 3.3. Rating System

**Decision:** Allow rating updates, calculate averages dynamically

#### Rationale:

- Users can modify their ratings as opinions change

- Real-time average calculation ensures data accuracy
- Simple aggregation suitable for small-scale operations

### 3.4. Recommendation Sharing

**Decision:** UUID-based public links with 7-day expiry

#### Rationale:

- UUIDs provide unpredictable, secure identifiers
- Public access eliminates authentication barriers for patients
- Expiry mechanism prevents indefinite access to potentially outdated recommendations

### 3.5. Product Integration

**Decision:** Fetch products synchronously from external API

#### Rationale:

- Real-time data ensures product information accuracy
- Graceful fallback for missing products
- Caching consideration for production environments

## 4. Input Validation and Business Rules

### 4.1. Validation Rules:

- Ratings: Integer range 1-5 with comprehensive error handling
- Reviews: 100-word limit with automatic word counting
- Authentication: Password hashing with SHA-256
- Recommendation expiry: Automatic 7-day TTL with server-side validation.

### 4.2. Business Logic:

- Doctors can only create recommendations (role-based authorization)
- Users can update existing ratings for the same doctor
- Public recommendation access without authentication requirements
- Expired recommendations return HTTP 410 Gone status

## 5. Assumptions and Limitations

### 5.1. Assumptions

1. **User Trust:** Users provide honest ratings and reviews
2. **Product API Reliability:** DummyJSON API maintains consistent availability
3. **Single-Doctor Rating:** Users rate each doctor only once (with updates allowed)
4. **English Content:** All text content is in English for word counting

### 5.2. Technical Limitations

1. **Scalability:** In-memory database unsuitable for production loads
2. **Concurrency:** No optimistic locking for concurrent rating updates
3. **Caching:** No caching layer for frequently accessed doctor lists
4. **Search:** Limited filtering capabilities (only by minimum rating)